

Package ‘ADMMsigma’

March 23, 2018

Type Package

Title Penalized Precision Matrix Estimation via ADMM

Version 1.0

Date 2018-02-23

Description This package estimates a penalized precision matrix via the alternating direction method of multipliers (ADMM) algorithm. It currently supports a general elastic-net penalty that allows for both ridge and lasso-type penalties as special cases.

URL <https://github.com/MGallow/ADMMsigma>

BugReports <https://github.com/MGallow/ADMMsigma/issues>

License GPL (>= 2)

ByteCompile TRUE

NeedsCompilation yes

Encoding UTF-8

LazyData true

RoxygenNote 6.0.1

Imports stats,
parallel,
foreach,
ggplot2

Depends Rcpp (>= 0.12.10),
doParallel

LinkingTo Rcpp,
RcppArmadillo

Suggests testthat

SystemRequirements GNU make

R topics documented:

ADMMsigma	2
plot.ADMMsigma	4
plot.RIDGESigma	5
RIDGESigma	5
Index	7

Description

Penalized precision matrix estimation using the ADMM algorithm. Consider the case where X_1, \dots, X_n are iid $N_p(\mu, \Sigma)$ and we are tasked with estimating the precision matrix, denoted $\Omega \equiv \Sigma^{-1}$. This function solves the following optimization problem:

$$\textbf{Objective: } \hat{\Omega}_\lambda = \arg \min_{\Omega \in S_+^p} \left\{ Tr(S\Omega) - \log \det(\Omega) + \lambda \left[\frac{1-\alpha}{2} \|\Omega\|_F^2 + \alpha \|\Omega\|_1 \right] \right\}$$

where $0 \leq \alpha \leq 1$, $\lambda > 0$, $\|\cdot\|_F^2$ is the Frobenius norm and we define $\|A\|_1 = \sum_{i,j} |A_{ij}|$. This elastic net penalty is identical to the penalty used in the popular penalized regression package `glmnet`. Clearly, when $\alpha = 0$ the elastic-net reduces to a ridge-type penalty and when $\alpha = 1$ this reduces to a lasso-type penalty.

Usage

```
ADMMsigma(X = NULL, S = NULL, lam = 10^seq(-5, 5, 0.5), alpha = seq(0,
  1, 0.1), diagonal = FALSE, rho = 2, mu = 10, tau1 = 2, tau2 = 2,
  crit = "ADMM", tol1 = 1e-04, tol2 = 1e-04, maxit = 1000, K = 5,
  cores = 1, quiet = TRUE)
```

Arguments

X	option to provide a nxp matrix. Each row corresponds to a single observation and each column contains n observations of a single feature/variable.
S	option to provide a pxp sample covariance matrix (denominator n). If argument is NULL and X is provided instead then S will be computed automatically.
lam	tuning parameter for elastic net penalty. Defaults to grid of values $10^{\text{seq}(-5, 5, 0.5)}$.
alpha	elastic net mixing parameter contained in [0, 1]. 0 = ridge, 1 = lasso. Defaults to grid of values $\text{seq}(-1, 1, 0.1)$.
diagonal	option to penalize the diagonal elements of the estimated precision matrix (Ω). Defaults to FALSE.
rho	initial step size for ADMM algorithm.
mu	factor for primal and residual norms in the ADMM algorithm. This will be used to adjust the step size rho after each iteration.
tau1	factor in which to increase step size rho
tau2	factor in which to decrease step size rho
crit	criterion for convergence (ADMM, grad, or loglik). If <code>crit != ADMM</code> then <code>tol1</code> will be used as the convergence tolerance. Default is ADMM.
tol1	absolute convergence tolerance. Defaults to 1e-4.
tol2	relative convergence tolerance. Defaults to 1e-4.
maxit	maximum number of iterations.
K	specify the number of folds for cross validation.
cores	option to run CV in parallel. Defaults to <code>cores = 1</code> .
quiet	specify whether the function returns progress of CV or not.

Details

For details on the implementation of 'ADMMsigma', see the vignette <https://mgallow.github.io/ADMMsigma/>.

Value

returns class object ADMMsigma which includes:

Iterations	number of iterations
Tuning	optimal tuning parameters (lam and alpha).
Lambdas	grid of lambda values for CV.
Alphas	grid of alpha values for CV.
maxit	maximum number of iterations.
Omega	estimated penalized precision matrix.
Sigma	estimated covariance matrix from the penalized precision matrix (inverse of Omega).
Gradient	gradient of optimization function (penalized gaussian likelihood).
CV.error	cross validation errors.

Author(s)

Matt Galloway <gall0441@umn.edu>

References

- For more information on the ADMM algorithm, see:
Boyd, Stephen, Neal Parikh, Eric Chu, Borja Peleato, Jonathan Eckstein, and others. 2011. 'Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers.' *Foundations and Trends in Machine Learning* 3 (1). Now Publishers, Inc.: 1-122.
https://web.stanford.edu/~boyd/papers/pdf/admm_distr_stats.pdf

See Also

[plot.ADMMsigma](#), [RIDGESigma](#)

Examples

```
# generate data from a dense matrix
# first compute covariance matrix
S = matrix(0.9, nrow = 5, ncol = 5)
diag(S) = 1

# generate 100 x 5 matrix with rows drawn from iid N_p(0, S)
Z = matrix(rnorm(100*5), nrow = 100, ncol = 5)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5)
S.sqrt = S.sqrt %*% t(out$vectors)
X = Z %*% S.sqrt

# elastic-net type penalty (use CV for optimal lambda and alpha)
ADMMsigma(X)
```

```
# ridge penalty (use CV for optimal lambda)
ADMMsigma(X, alpha = 0)

# lasso penalty (lam = 0.1)
ADMMsigma(X, lam = 0.1, alpha = 1)

# produce CV heat map for ADMMsigma
plot(ADMMsigma(X))
```

plot.ADMMSigma	<i>Plot ADMMsigma object</i>
----------------	------------------------------

Description

produces a heat plot for the cross validation errors, if available.

Usage

```
## S3 method for class 'ADMMsigma'
plot(x, footnote = TRUE, ...)
```

Arguments

x	class object ADMMsigma.
footnote	option to print footnote of optimal values.
...	additional arguments.

Examples

```
# generate data from a dense matrix
# first compute covariance matrix
S = matrix(0.9, nrow = 5, ncol = 5)
diag(S) = 1

# generate 100 x 5 matrix with rows drawn from iid N_p(0, S)
Z = matrix(rnorm(100*5), nrow = 100, ncol = 5)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5)
S.sqrt = S.sqrt %*% t(out$vectors)
X = Z %*% S.sqrt

# produce CV heat map for ADMMsigma
plot(ADMMsigma(X))
```

plot.RIDGESigma	<i>Plot RIDGESigma object</i>
-----------------	-------------------------------

Description

produces a heat plot for the cross validation errors, if available.

Usage

```
## S3 method for class 'RIDGESigma'
plot(x, footnote = TRUE, ...)
```

Arguments

x	class object RIDGESigma
footnote	option to print footnote of optimal values.
...	additional arguments.

Examples

```
# generate data from a dense matrix
# first compute covariance matrix
S = matrix(0.9, nrow = 5, ncol = 5)
diag(S) = 1

# generate 100 x 5 matrix with rows drawn from iid N_p(0, S)
Z = matrix(rnorm(100*5), nrow = 100, ncol = 5)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5)
S.sqrt = S.sqrt %*% t(out$vectors)
X = Z %*% S.sqrt

# produce CV heat map for RIDGESigma
plot(RIDGESigma(X, lam = 10^seq(-8, 8, 0.01)))
```

RIDGESigma	<i>Ridge penalized precision matrix estimation</i>
------------	--

Description

Ridge penalized matrix estimation via closed-form solution. If you are only interested in the ridge penalty, this function will be faster and provide a more precise estimate than using ADMMsigma. Consider the case where X_1, \dots, X_n are iid $N_p(\mu, \Sigma)$ and we are tasked with estimating the precision matrix, denoted $\Omega \equiv \Sigma^{-1}$. This function solves the following optimization problem:

Objective: $\hat{\Omega}_\lambda = \arg \min_{\Omega \in S_+^p} \left\{ Tr(S\Omega) - \log \det(\Omega) + \frac{\lambda}{2} \|\Omega\|_F^2 \right\}$

where $\lambda > 0$ and $\|\cdot\|_F^2$ is the Frobenius norm.

Usage

```
RIDGEsigma(X = NULL, S = NULL, lam = 10^seq(-5, 5, 0.5), K = 3,
  quiet = TRUE)
```

Arguments

X	option to provide a nxp matrix. Each row corresponds to a single observation and each column contains n observations of a single feature/variable.
S	option to provide a pxp sample covariance matrix (denominator n). If argument is NULL and X is provided instead then S will be computed automatically.
lam	tuning parameter for ridge penalty. Defaults to grid of values $10^{\text{seq}(-5, 5, 0.5)}$.
K	specify the number of folds for cross validation.
quiet	specify whether the function returns progress of CV or not.

Value

returns class object RIDGESigma which includes:

Lambda	optimal tuning parameter.
Lambdas	grid of lambda values for CV.
Omega	estimated penalized precision matrix.
Sigma	estimated covariance matrix from the penalized precision matrix (inverse of Omega).
Gradient	gradient of optimization function (penalized gaussian likelihood).
CV.error	cross validation errors.

Author(s)

Matt Galloway <gall0441@umn.edu>

See Also

[plot.RIDGESigma](#), [ADMMsigma](#)

Examples

```
# generate data from a dense matrix
# first compute covariance matrix
S = matrix(0.9, nrow = 5, ncol = 5)
diag(S) = 1

# generate 100 x 5 matrix with rows drawn from iid N_p(0, S)
Z = matrix(rnorm(100*5), nrow = 100, ncol = 5)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5)
S.sqrt = S.sqrt %*% t(out$vectors)
X = Z %*% S.sqrt

# ridge penalty no ADMM
RIDGEsigma(X, lam = 10^seq(-8, 8, 0.01))

# produce CV heat map for RIDGESigma
plot(RIDGEsigma(X, lam = 10^seq(-8, 8, 0.01)))
```

Index

ADMMsigma, [2](#), [6](#)

plot.ADMMsigma, [3](#), [4](#)

plot.RIDGESigma, [5](#), [6](#)

RIDGESigma, [3](#), [5](#)