Regularized Precision Matrix Estimation via ADMM

Matt Galloway

March 17, 2018

Abstract

`ADMMsigma` is an R package that estimates a penalized precision matrix via the alternating direction method of multipliers (ADMM) algorithm. This report will provide a brief overview of the algorithm and detail how it can be utilized to estimate precision matrices of joint normal distributions. In addition, examples and simulation results will be provided for `ADMMsigma`.

Contents

1   Introduction

Suppose we want to minimize $f(x) + g(z)$ subject to the constraint that $Ax + Bz = c$. For now, we will take $x \in \mathbb{R}^n, z \in \mathbb{R}^m, A \in \mathbb{R}^{p \times m}, B \in \mathbb{R}^{p \times m}, c \in \mathbb{R}^p$ – though we will later consider cases where x and z are matrices. The *augmented lagrangian* is constructed as follows:

$$L_\rho(x, z, y) = f(x) + g(z) + y^T(Ax + Bz - c) + \frac{\rho}{2}\|Ax + Bz - c\|_2^2$$

where $y \in \mathbb{R}^p$ is the lagrange multiplier. The optimal value is

$$p^* = inf\{f(x) + g(z)|Ax + Bz = c\}$$

Clearly, the minimization problem under the augmented lagrangian (RE-WORK) is equivalent to that of the usual lagrangian since any feasible point (x, z) satisfies the constraint $\rho\|Ax + Bz - c\|_2^2/2 = 0$.

The ADMM algorithm consists of the following repeated iterations:

1

$$x^{k+1} := \arg\min_x L_\rho(x, z^k, y^k) \tag{1}$$

$$z^{k+1} := \arg\min_z L_\rho(z^{k+1}, z, y^k) \tag{2}$$

$$y^{k+1} := y^k + \rho(Ax^{k+1} + Bz^{k+1} - c) \tag{3}$$

A more complete introduction to the algorithm – specifically how it arose out of *dual ascent* and *method of multipliers* – can be found in Boyd et al. (2011).

2 Regularized Precision Matrix Estimation

We now consider the case where $X_1, ..., X_n$ are iid $N_p(\mu, \Sigma)$ and we are tasked with estimating the precision matrix, denoted $\Omega \equiv \Sigma^{-1}$. The maximum likelihood estimator for Ω is

$$\hat{\Omega} = \arg\min_{\Omega \in S_+^p} \{Tr(S\Omega) - \log\det(\Omega)\}$$

where $S = \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T/n$. It is straight forward to show that when the solution exists, $\hat{\Omega} = S^{-1}$.

We can further construct a penalized likelihood estimator by adding a penalty term, $P_\lambda(\Omega)$, to the likelihood:

$$\hat{\Omega}_\lambda = \arg\min_{\Omega \in S_+^p} \{Tr(S\Omega) - \log\det(\Omega) + P_\lambda(\Omega)\}$$

Throughout the rest of this document we will take $P_\lambda(\Omega)$ to be $P_\lambda(\Omega) = \lambda\left[\frac{1-\alpha}{2}\|\Omega\|_F^2 + \alpha\|\Omega\|_1\right]$ so that the full penalized likelihood is as follows:

$$\hat{\Omega}_\lambda = \arg\min_{\Omega \in S_+^p} \left\{Tr(S\Omega) - \log\det(\Omega) + \lambda\left[\frac{1-\alpha}{2}\|\Omega\|_F^2 + \alpha\|\Omega\|_1\right]\right\}$$

where $0 \le \alpha \le 1$, $\lambda > 0$, $0 < \eta < 2$, $\|\cdot\|_F^2$ is the Frobenius norm and we define $\|A\|_1 = \sum_{i,j} |A_{ij}|$. This penalty is closely related to the elastic-net penalty explored by Hui Zou and Trevor Hastie (Zou and Hastie 2005). Clearly, when $\alpha = 0$ this reduces to a ridge-type penalty and when $\alpha = 1$ this reduces to a lasso-type penalty.

By letting f be equal to the non-penalized likelihood and g equal to $P_\lambda(\Omega)$, our goal is to minimize the full augmented lagrangian where the constraint is that $\Omega - Z$ is equal to zero:

$$L_\rho(\Omega, Z, \Lambda) = f(\Omega) + g(Z) + Tr[\Lambda(\Omega - Z)] + \frac{\rho}{2}\|\Omega - Z\|_F^2$$

The ADMM algorithm for regularized precision matrix estimation is

$$\Omega^{k+1} = \arg\min_{S\Omega}\left\{Tr(\Omega) - \log\det(\Omega) + Tr[\Lambda^k(\Omega - Z^k)] + \frac{\rho}{2}\|\Omega - Z^k\|_F^2\right\} \tag{4}$$

$$Z^{k+1} = \arg\min_Z\left\{\lambda\left[\frac{1-\alpha}{2}\|Z\|_F^2 + \alpha\|Z\|_1\right] + Tr[\Lambda^k(\Omega^{k+1} - Z)] + \frac{\rho}{2}\|\Omega^{k+1} - Z\|_F^2\right\} \tag{5}$$

$$\Lambda^{k+1} = \Lambda^k + \rho(\Omega^{k+1} - Z^{k+1}) \tag{6}$$

2.1 Condensed-Form ADMM

An alternate form of the ADMM algorithm can constructed by scaling the dual variable. Let us define $R^k = \Omega - Z^k$ and $U^k = \Lambda^k/\rho$. Then

$$Tr\left[\Lambda^k\left(\Omega - Z^k\right)\right] + \frac{\rho}{2}\left\|\Omega - Z^k\right\|_F^2 = Tr\left[\Lambda^k R^k\right] + \frac{\rho}{2}\left\|R^k\right\|_F^2$$
$$= \frac{\rho}{2}\left\|R^k + \Lambda^k/\rho\right\|_F^2 - \frac{\rho}{2}\left\|\Lambda^k/\rho\right\|_F^2$$
$$= \frac{\rho}{2}\left\|R^k + U^k\right\|_F^2 - \frac{\rho}{2}\left\|U^k\right\|_F^2$$

The condensed-form can now be written as follows:

$$\Omega^{k+1} = \arg\min_{\Omega}\left\{Tr\left(\Omega\right) - \log\det\left(\Omega\right) + \frac{\rho}{2}\left\|\Omega - Z^k + U^k\right\|_F^2\right\} \tag{7}$$

$$Z^{k+1} = \arg\min_{Z}\left\{\lambda\left[\frac{1-\alpha}{2}\left\|Z\right\|_F^2 + \alpha\left\|Z\right\|_1\right] + \frac{\rho}{2}\left\|\Omega^{k+1} - Z + U^k\right\|_F^2\right\} \tag{8}$$

$$U^{k+1} = U^k + \Omega^{k+1} - Z^{k+1} \tag{9}$$

More generally (in vector form),

$$x^{k+1} := \arg\min_{x}\left\{f(x) + \frac{\rho}{2}\left\|Ax + Bz^k - c + u^k\right\|_2^2\right\} \tag{10}$$

$$z^{k+1} := \arg\min_{z}\left\{g(z) + \frac{\rho}{2}\left\|Ax^{k+1} + Bz - c + u^k\right\|_2^2\right\} \tag{11}$$

$$u^{k+1} := u^k + Ax^{k+1} + Bz^{k+1} - c \tag{12}$$

Note that there are limitations to using this method. For instance, because the dual variable is scaled by ρ (the step size), this form limits one to using a constant step size (without making further adjustments to U^k) – a limitation that could prolong the convergence rate.

2.2 Algorithm

$$\Omega^{k+1} = \arg\min_{\Omega}\left\{Tr\left(\Omega\right) - \log\det\left(\Omega\right) + \frac{\rho}{2}\left\|\Omega - Z^k + U^k\right\|_F^2\right\}$$

$$Z^{k+1} = \arg\min_{Z}\left\{\lambda\left[\frac{1-\alpha}{2}\left\|Z\right\|_F^2 + \alpha\left\|Z\right\|_1\right] + \frac{\rho}{2}\left\|\Omega^{k+1} - Z + U^k\right\|_F^2\right\}$$

$$U^{k+1} = U^k + \Omega^{k+1} - Z^{k+1}$$

1. Decompose $S + \rho(U^k - Z^k) = VQV^T$.

$$\Omega^{k+1} = \frac{1}{2\rho}V\left[-Q + \left(Q^2 + 4\rho I_p\right)^{1/2}\right]V^T$$

2. Elementwise soft-thresholding for all $i = 1, ..., p$ and $j = 1, ..., p$.

$$Z_{ij}^{k+1} = \frac{1}{\lambda(1 - \alpha) + \rho} sign\left(\Omega_{ij}^{k+1} + U_{ij}^k\right)\left(\rho\left|\Omega_{ij}^{k+1} + U_{ij}^k\right| - \lambda\eta\alpha\right)_+$$

$$= \frac{1}{\lambda(1 - \alpha) + \rho} Soft\left(\rho(\Omega_{ij}^{k+1} + U_{ij}^k), \lambda\eta\alpha\right)$$

3. Update U.

$$U^{k+1} = U^k + \Omega^{k+1} - Z^{k+1}$$

2.2.1 Proof of (1):

(Work in progress.)

Code snippet:

Note this is not the actual code. The real code is written in c++.

```r
# ridge penalized precision matrix
# function
RIDGEsigma = function(S, lam) {

    # dimensions
    p = dim(S)[1]

    # gather eigen values of S (spectral
    # decomposition)
    e.out = eigen(S, symmetric = TRUE)

    # augment eigen values for omega hat
    new.evs = (-e.out$val + sqrt(e.out$val^2 +
        4 * lam))/(2 * lam)

    # compute omega hat for lambda (zero
    # gradient equation)
    omega = tcrossprod(e.out$vec * rep(new.evs,
        each = p), e.out$vec)

    # compute gradient
    grad = S - qr.solve(omega) + lam * omega

    return(list(omega = omega, gradient = grad))
}
```

4

2.2.2 Proof of (2)

(Work in progress.)

Code snippet:

Note this is not the actual code. The real code is written in c++.

```
# ADMMsigma function
ADMMsigma = function(X = NULL, S = NULL,
    lam, alpha = 1, rho = 2, mu = 10, tau1 = 2,
    tau2 = 2, tol1 = 1e-04, tol2 = 1e-04,
    maxit = 1000) {

    # compute sample covariance matrix, if
    # necessary
    if (is.null(S)) {

        # covariance matrix
        n = dim(X)[1]
        S = (n - 1)/n * cov(X)

    }

    # allocate memory
    p = dim(S)[1]
    criterion = TRUE
    iter = lik = s = r = eps1 = eps2 = 0
    new.Z = Y = Omega = matrix(0, nrow = p,
        ncol = p)

    # loop until convergence
    while (criterion && (iter <= maxit)) {

        # ridge equation (1) gather eigen values
        # (spectral decomposition)
        Z = new.Z
        Omega = sigma_ridge(S + Y - rho *
            Z, lam = rho)$omega

        # penalty equation (2) soft-thresholding
        new.Z = soft(Y + rho * Omega, lam *
            alpha)/(lam * (1 - alpha) + rho)

        # update U (3)
        Y = Y + rho * (Omega - new.Z)

        # calculate new rho
        s = sqrt(sum((rho * (new.Z - Z))^2))
        r = sqrt(sum((Omega - new.Z)^2))
        rho = rho * (tau1 * (r > mu * s) +
            (s > mu * r)/tau2 + (s/mu <=
```

```
            r & r <= mu * s))
        iter = iter + 1

        # stopping criterion
        eps1 = p * tol1 + tol2 * max(sqrt(sum(Omega^2)),
            sqrt(sum(new.Z^2)))
        eps2 = p * tol1 + tol2 * sqrt(sum(Y^2))
        criterion = (r >= eps1 || s >= eps2)

    }
    return(list(Iterations = iter, Omega = Omega))
}
```

3   R Package

3.1   Installation

```
# The easiest way to install is from the
# development version from GitHub:
# install.packages('devtools')
devtools::install_github("MGallow/ADMMsigma")
```

If there are any issues/bugs, please let me know: github. You can also contact me via my website. Pull requests are welcome!

3.2   Usage

```
library(ADMMsigma)

# generate data from a dense matrix first
# compute covariance matrix
S = matrix(0, nrow = 5, ncol = 5)

for (i in 1:5) {
    for (j in 1:5) {
        S[i, j] = 0.9^(i != j)
    }
}
```

```
# generate 100 x 5 matrix with rows drawn
# from iid N_p(0, S)
Z = matrix(rnorm(100 * 5), nrow = 100, ncol = 5)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5) %*%
    t(out$vectors)
X = Z %*% S.sqrt


# elastic-net type penalty (use CV for
# optimal lambda and alpha)
ADMMsigma(X)
```

```
##
## Iterations:
## [1] 35
##
## Tuning parameters:
##       log10(lam)  alpha
## [1,]        -2.5    0.6
##
## Omega:
##         [,1]     [,2]     [,3]     [,4]     [,5]
## [1,]  7.50103 -1.71813 -2.17055 -1.21055 -1.95217
## [2,] -1.71813  6.88957 -2.51084 -1.42759 -1.14920
## [3,] -2.17055 -2.51084  8.17508 -1.79449 -1.38145
## [4,] -1.21055 -1.42759 -1.79449  6.25581 -1.97081
## [5,] -1.95217 -1.14920 -1.38145 -1.97081  6.53472
```
```
# ridge penalty (use CV for optimal
# lambda)
ADMMsigma(X, alpha = 0)
```

```
##
## Iterations:
## [1] 39
##
## Tuning parameters:
##       log10(lam)  alpha
## [1,]           -3      0
##
## Omega:
##         [,1]     [,2]     [,3]     [,4]     [,5]
## [1,]  7.92134 -1.80920 -2.32991 -1.24122 -2.07704
## [2,] -1.80920  7.24953 -2.70397 -1.48210 -1.17634
## [3,] -2.32991 -2.70397  8.69320 -1.89807 -1.43455
## [4,] -1.24122 -1.48210 -1.89807  6.54511 -2.08445
## [5,] -2.07704 -1.17634 -1.43455 -2.08445  6.84977
```
```
# lasso penalty (lam = 0.1)
ADMMsigma(X, lam = 0.1, alpha = 1)
```

```
##
## Iterations:
## [1] 10
```

```
## 
## Tuning parameters:
##       log10(lam)  alpha
## [1,]          -1      1
## 
## Omega:
##           [,1]     [,2]     [,3]     [,4]     [,5]
## [1,]   2.80422 -0.62667 -0.67987 -0.55554 -0.66056
## [2,]  -0.62667  2.66766 -0.75422 -0.60328 -0.54369
## [3,]  -0.67987 -0.75422  2.88706 -0.65217 -0.57631
## [4,]  -0.55554 -0.60328 -0.65217  2.53979 -0.69921
## [5,]  -0.66056 -0.54369 -0.57631 -0.69921  2.60963
```
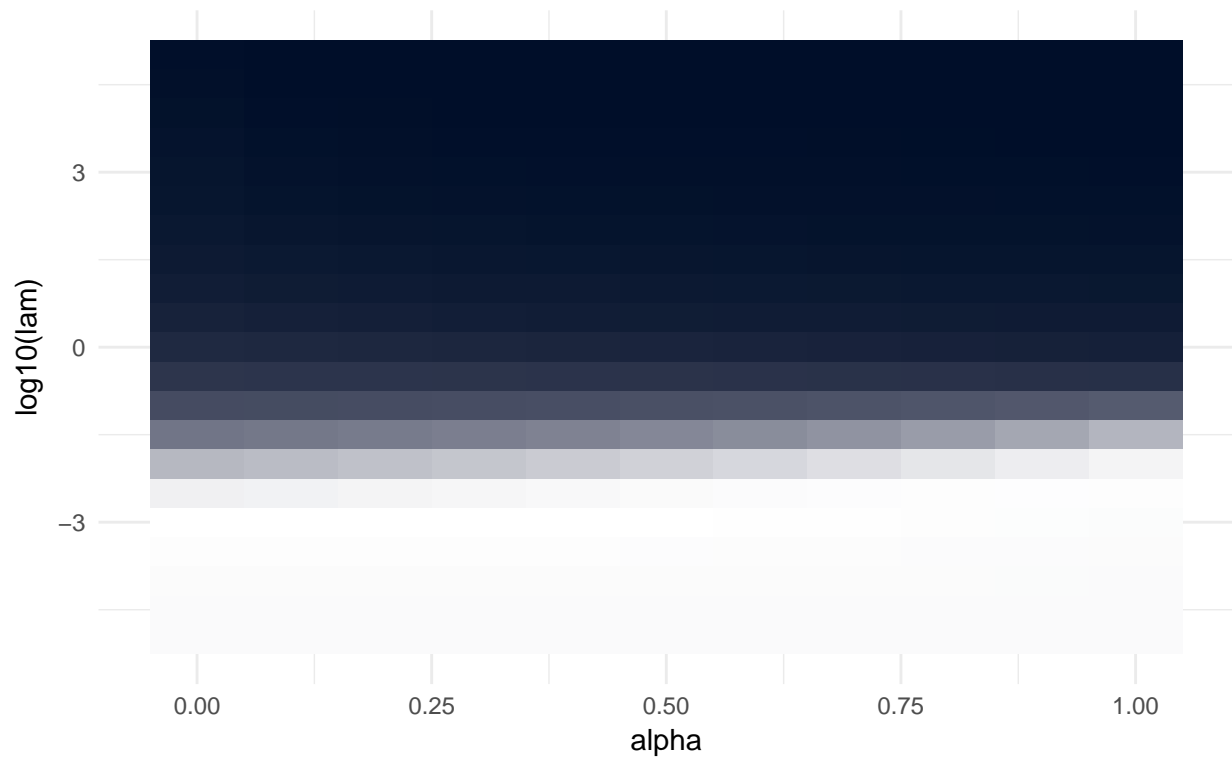
```r
# ridge penalty no ADMM
RIDGEsigma(X, lam = 10^seq(-8, 8, 0.01))
```

```
## 
## Tuning parameter:
##         lam  log10(lam)
## [1,]  0.001        -3.1
## 
## Omega:
##           [,1]     [,2]     [,3]     [,4]     [,5]
## [1,]   8.08365 -1.83697 -2.40822 -1.24051 -2.12903
## [2,]  -1.83697  7.38256 -2.79786 -1.49474 -1.17783
## [3,]  -2.40822 -2.79786  8.92397 -1.94165 -1.44583
## [4,]  -1.24051 -1.49474 -1.94165  6.63854 -2.12649
## [5,]  -2.12903 -1.17783 -1.44583 -2.12649  6.95522
```

```r
# produce CV heat map for ADMMsigma
ADMMsigma(X) %>% plot
```
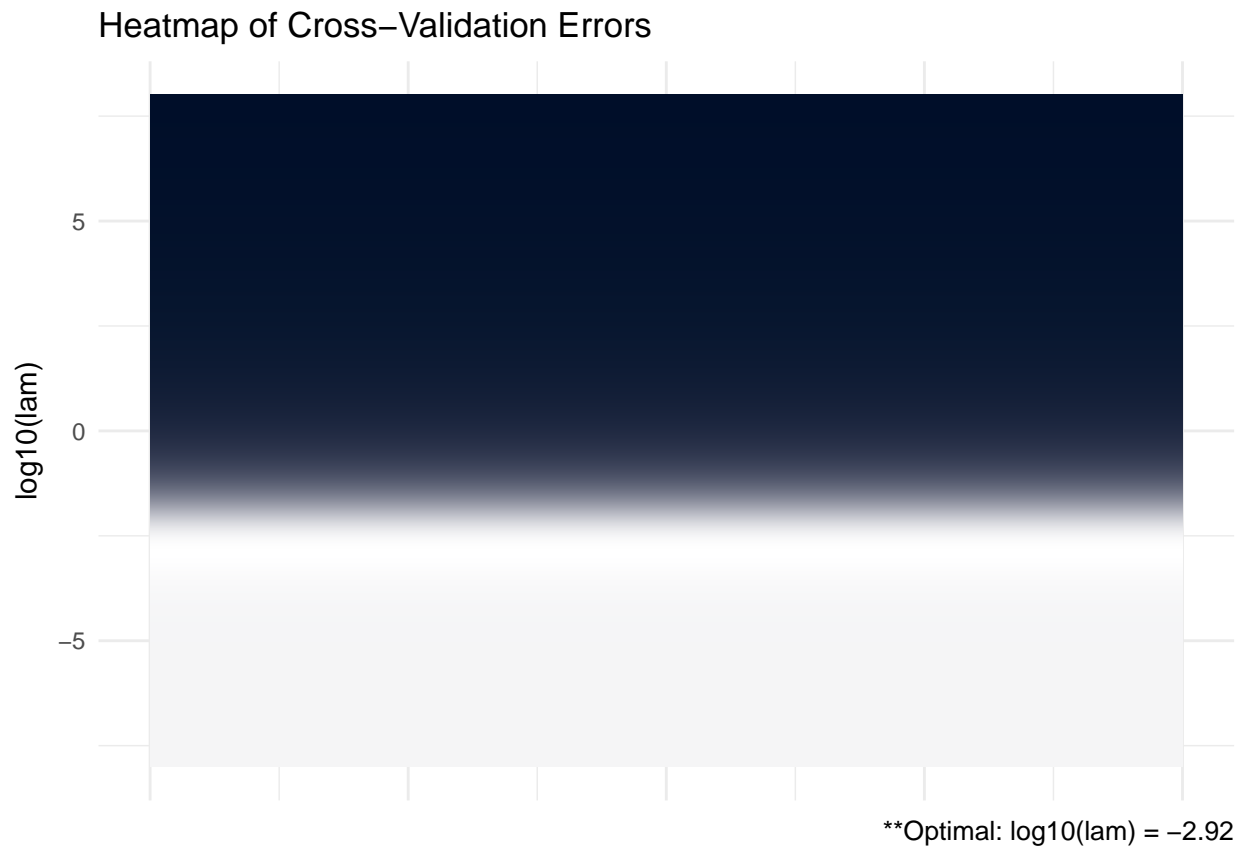
Heatmap of Cross−Validation Errors

**Optimal: log10(lam) = −3, alpha = 0.2

```
# produce CV heat map for RIDGEsigma
RIDGEsigma(X, lam = 10^seq(-8, 8, 0.01)) %>%
    plot
```

Heatmap of Cross−Validation Errors

**Optimal: log10(lam) = −2.92

3.3 Benchmark

3.3.1 Computer Specs:

- MacBook Pro (Late 2016)
- Processor: 2.9 GHz Intel Core i5
- Memory: 8GB 2133 MHz
- Graphics: Intel Iris Graphics 550

```
# generate data from tri-diagonal
# (sparse) matrix compute covariance
# matrix (can confirm inverse is
# tri-diagonal)
S = matrix(0, nrow = 100, ncol = 100)

for (i in 1:100) {
    for (j in 1:100) {
        S[i, j] = 0.7^(abs(i - j))
    }
}

# generate 1000 x 100 matrix with rows
# drawn from iid N_p(0, S)
Z = matrix(rnorm(1000 * 100), nrow = 1000,
```

```
    ncol = 100)
out = eigen(S, symmetric = TRUE)
S.sqrt = out$vectors %*% diag(out$values^0.5) %*%
    t(out$vectors)
X = Z %*% S.sqrt


# glasso (for comparison)
microbenchmark(glasso(s = S, rho = 0.1))
```

```
## Unit: milliseconds
##                      expr      min       lq     mean   median       uq
##  glasso(s = S, rho = 0.1) 49.93206 53.11861 57.87519 55.44664 59.30303
##      max neval
##  89.96595   100
```

```
# benchmark ADMMsigma - default tolerance
microbenchmark(ADMMsigma(S = S, lam = 0.1,
    alpha = 1, tol1 = 1e-04, tol2 = 1e-04))
```

```
## Unit: milliseconds
##                                                              expr
##  ADMMsigma(S = S, lam = 0.1, alpha = 1, tol1 = 1e-04, tol2 = 1e-04)
##      min       lq     mean   median       uq      max neval
##  39.99944 43.82541 51.62732 46.94982 53.30376 228.9022   100
```

```
# benchmark ADMMsigma - tolerance 1e-8
microbenchmark(ADMMsigma(S = S, lam = 0.1,
    alpha = 1, tol1 = 1e-08, tol2 = 1e-08))
```

```
## Unit: milliseconds
##                                                              expr
##  ADMMsigma(S = S, lam = 0.1, alpha = 1, tol1 = 1e-08, tol2 = 1e-08)
##      min       lq     mean   median       uq      max neval
##  189.2626 203.0107 221.2471 211.5412 230.3703 361.4178   100
```

```
# benchmark ADMMsigma CV - default
# parameter grid
microbenchmark(ADMMsigma(X), times = 5)
```

```
## Unit: seconds
##          expr      min       lq     mean   median       uq      max neval
##  ADMMsigma(X) 55.62628 56.18889 58.61224  56.4466 59.59131 65.20811     5
```

```
# benchmark ADMMsigma parallel CV
microbenchmark(ADMMsigma(X, parallel = TRUE),
    times = 5)
```

```
## Unit: seconds
##                          expr      min       lq     mean   median
##  ADMMsigma(X, parallel = TRUE) 30.97527 31.30452 31.75878 32.06775
##      uq      max neval
##  32.22213 32.22423     5
```

```
# benchmark ADMMsigma CV - likelihood
# convergence criteria
microbenchmark(ADMMsigma(X, crit = "loglik"),
```

```
    times = 5)
```

```
## Unit: seconds
##                               expr      min       lq     mean   median
##  ADMMsigma(X, crit = "loglik") 80.37973 82.30604 83.37748 82.87069
##        uq     max neval
##  85.60625 85.7247     5
```

References

Boyd, Stephen, Neal Parikh, Eric Chu, Borja Peleato, Jonathan Eckstein, and others. 2011. "Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers." *Foundations and Trends in Machine Learning* 3 (1). Now Publishers, Inc.: 1–122.

Zou, Hui, and Trevor Hastie. 2005. "Regularization and Variable Selection via the Elastic Net." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 67 (2). Wiley Online Library: 301–20.